

# Lab Session 1

## Introduction

In this Lab exercise you'll be introduced to:

1. A small warmup exercise with numpy, Matplotlib.
2. Best practise to documenting your code.
3. Bandit Algorithms

## Instructions

In this lab session you will have to implement Bandit algorithms, Bandits (aka slot machines) and the learning loop of the agents. The code is divided in 4 modules, and you will have to complete the code in the 3 first ones:

### Module 1: launcher.py

This module controls the hyperparameters, main loops, and plotting.

#### General functioning

You are given two possibilities for launch, using the `launch\_type` variable:

1. *multiple\_agents*: comparing multiple agents on the k-armed bandit problem. Fill the list of agent classes you want to compare, `agents`; you can change the labels and plot file names.
2. *spectrum*: comparing values of a hyperparameter for an algorithm on the k-armed bandit problem. The `spectrum` variable should look like a list `['hparam\_name', [list of values the hparam should take]]`; you can change the labels and plot file names.

**Note:-** you cannot run the code before you implement at least

- \* Gaussian\_Bandit
- \* KBandit
- \* run\_bandit

#### TODO

You only have to fill in the *run\_bandit* function. The details of the function are in its documentation.

### Module 2: bandit.py

This module contains the *Slot Machines* of the problem.

A Slot Machine is actually called a "bandit" because it steals money from the gambler's pockets, apparently.

#### General functioning

The *Bandit* class is implemented: it is the abstraction of a machine that you can simply pull, and get a reward.

All the subsequent classes should be child classes of Bandit.

## TODO

You need to implement the following classes:

- \* *Gaussian\_Bandit* - the machine's reward distribution is a Gaussian  $N(m,1)$  around a fixed mean  $m \sim N(1,0)$
- \* *Gaussian\_Bandit\_NonStat* - a Gaussian\_Bandit whose mean moves  $\sim N(0, \text{std}=0.01)$  after each pull (of any other arm also)
- \* *KBandit* - a k-armed bandit, i.e. a set of k Gaussian\_Bandits
- \* *KBandit\_NonStat* - a KBandit, but with Non Stationary Bandits, i.e. the distributions move every pull.

All the classes should be children (or grand-) of the abstract Bandit class, to keep a clean **object-oriented code**.

Once the KBandit and run\_bandit class and function have been implemented, you can try to run the Random\_Agent (given as an example) by selecting ``launch_type='mutiple_agents'`` and only un-commenting the ``Random_Agent`` from the list of ``agents``. You can set as output file\_name, for example `'random_agent'`.

Do the results make sense?

## Module 3: agents.py

This module contains the agents, or algorithms, that are meant to solve the Bandit problems implemented above.

### General functioning

The *Bandit\_Agent* abstract class is implemented and contains the formalism necessary to run an agent. Such an agent should be able to act, learn and reset itself.

The *Random\_Agent* class was implemented to show an example, but note how it does not need to implement the ``learn`` method, by definition.

## TODO

As soon as a class is implemented, you can test it by comparing it to just the Random Agent (or another class you implemented) in the ``agents`` list.

The agent classes you will have to implement are:

- \* *EpsGreedy\_SampleAverage* - estimating Q values with sample averages and selecting actions epsilon-greedily
- \* *EpsGreedy* - fixed learning rate instead of Sample Averages
- \* *OptimisticGreedy* - high initial Q values; full greedy action selection
- \* *UCB* - selecting action greedily according to  $Q + U$  where U is a count-based bonus on rarely tried action
- \* *Gradient* - keeping up advantages vector H for each action, action selection samples from  $\text{softmax}(H)$

## Module 4: utils.py

Contains the plotting functions, some useful functions like the softmax or `my_random_choice`.

You **don't need to do anything** here. Enjoy!

## Submitting

Please record your plots in the folder of the same name (the `file\_name` variable controls this in addition to the plot name). We ask that you at least provide a comparison of all agents (both action and average reward) (file\_name='plots/agent\_comparison'). Once done, please zip your lab\_session1 folder and submit it on blackboard.